

# Software Architecture & Databases

Reicman University

# Contact Me

amitm

@

[at-bay.com](mailto:amitm@at-bay.com)

**Feel free to drop me your comments, suggestions, questions**

# Nice to Meet You!

- Life: 45 yo, married + 3 kids, live in Shoham, Yoga.
- Current: 2 yrs @ [At-Bay](#) as a Software Architect (...what's that?)
- Overall: ~15 yrs in the industry, mostly in startups, but fallen a couple of times to the corporate hands
- Fun time 1: Does Tech Due Diligences for a couple of VCs
- Fun time 2: Contributor to the [Dapr](#) open source project
- Fun time 3: this thing

# What We'll Talk About

- Software architecture - theory
- Where theory meets reality
- The oh-so-boring SaaS company
- The *relational database* as the greatest invention in industrial software engineering
- Where the relational database *fails*
- Q & A

# Software Architecture

## Wikipedia

- "Software architecture refers to the fundamental **structures** of a software **system** and the **discipline** of creating such structures and systems."
- "Each **structure** comprises software **elements**, **relations** among them, and **properties** of both elements and relations."
- system:
  - structures:
    - $\text{element} \leftarrow\!\!\rightarrow \text{relation} \rightarrow\!\!\rightarrow \text{element}$

# Software Architecture

## Somebody Needs to Choose... ("System Analysis")

- "Software architecture is about making fundamental structural **choices** that are **costly to change once implemented**" (wikipedia, this time getting it right)
- What's costly to change (over time):
  - Programming Language (changing/adding mostly breaks common tools) --> "element"
  - **Data Model** (breaks everything if done wrong) --> "property of an element"
  - Runtime Environment (on-prem vs cloud x/y/z) --> "structure+system"
  - Contracts (APIs) --> "relations"

# Software Architecture

## Conceptualization & Creation of New Things ("System Design")

- "Fighting" the windmills of **complexity** as software obeys the 2nd law of thermodynamics (complexity never decreases)
- But what is complex? ("This part of the codebase doesn't *feel* right!")
- When do you surrender to the evil called "management"?
- The humanity! (You care about people not losing their minds)



# Software Architecture

## Reviewing Things ("Evaluation")

- High Level Design Review - system
- Code Reviews - quality
- Security Reviews - security safety
- Test Review - code safety
- Data Modeling Review - correct abstraction/database performance
- Data Integration Review - none applicative systems
- DevOps Review - environment/deployment
- SRE Review - gauges

Make sure stuff aligns with your company *vision* and *business status* on how software is being made



# Software Architecture

## Redesign Bad Stuff Other People Did ("Evolution")

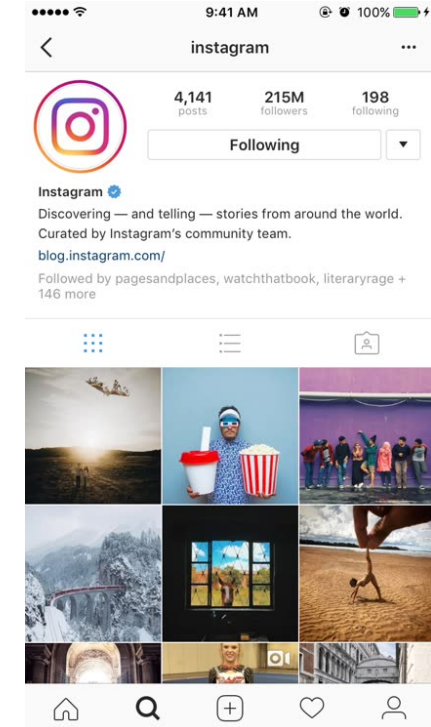
- Fun/easy stuff is to do *NEW* things
- NotSoFun/hardest stuff is to evolve *OLD* things
- Read about: "evolutionary architecture"

# Ope There Goes Gravity (E. minem)

- The ideal picture is a fantasy
- You have to **surrender** to accepting these:
  - You and your work are **not eternal**
  - There is no "generic" best architecture
  - Even the "best" code is eventually **thrown** to the recycle bean
  - You are working for a company that in its essence is a machine made to make **profit**
  - At the broader perspective, *individual contributions*, are **irrelevant**. Collaboration is key.
- None of the above is in contradiction to software development being a beautiful human mental act

# The Oh-So-Boring SaaS Company

- Capture *Data*
- Do something over *Data*
- Arrange the *Data* nicely
- Ask for money from customers/advertisers



# The Data Model

## Data Model = Database Tables

- It is how you model the business (entities)
- **It is the single most important element in the system design of a SaaS**
- It is the only effort worth investing BEFORE writing a single line of code (atypical for a startup to do)
- Still, it is ever **evolving** with the business needs, but once in use, *it is hard to change*
- Things that aren't persisted are potentially lost, so the use of in-memory data structures (with all due respect to BigO) is just temporary and the **significant stateful operations are done over a database**

# Data Model - Implementation


Usually, we're talking about a *Relational Database*

Rank			DBMS	Database Model	Score		
May 2022	Apr 2022	May 2021			May 2022	Apr 2022	May 2021
1.	1.	1.	Oracle +	Relational, Multi-model i	1262.82	+8.00	-7.12
2.	2.	2.	MySQL +	Relational, Multi-model i	1202.10	-2.06	-34.28
3.	3.	3.	Microsoft SQL Server +	Relational, Multi-model i	941.20	+2.74	-51.46
4.	4.	4.	PostgreSQL +	Relational, Multi-model i	615.29	+0.83	+56.04
5.	5.	5.	IBM Db2	Relational, Multi-model i	160.32	-0.13	-6.34
6.	6.	↑ 7.	Microsoft Access	Relational	143.44	+0.66	+28.04
7.	7.	↓ 6.	SQLite +	Relational	134.73	+1.94	+8.04
8.	8.	8.	MariaDB +	Relational, Multi-model i	111.13	+0.81	+14.44
9.	9.	↑ 16.	Snowflake +	Relational	93.51	+4.06	+63.46
10.	10.	10.	Microsoft Azure SQL Database	Relational, Multi-model i	85.33	-0.45	+14.88

# Relational Databases

The founding father of the relational model is  
*Edgar F. Codd*

- Based on his work published in 1970 (yes, 52 yo technology) while he was working for IBM
- He won the Turing Award in 1981 for this work
- He applied Relational Algebra and proposed such an algebra as a basis for database query languages
- Five primitive operators: selection, projection, Cartesian product (also called cross join), set union, set difference.



# Relational Databases - Basics

- Relational databases are all about **tables**
- Tables are able to *relate* to one another

Positions

id (pk)	name	start_date	end_date	employee_id (fk)
a21sazn	Payoneer	1/2/2019	1/6/2020	t2a9bc
b3at78	Payoneer	1/4/2019	1/10/2020	dlli89
k8aa6d	At-Bay	7/6/2020	null	t2a9bc

People

id (pk)	first	last	yob
t2a9bc	amit	mor	1977
dlli89	nir	pinchas	1981

Relation!

# Relational Databases - Basics 2

- Use a **declarative** language to apply operations over the logical representation and don't mind the physical aspects
- Ability to create meaningful information by **joining** of tables
- To ensure that data is always accurate and accessible, relational databases follow certain integrity rules (**A.C.I.D**)
- Relational databases are **transactional**—they guarantee the state of the entire system is consistent at any moment
- The relational model means that the *logical data structures*—the data tables, views, and indexes—are separate from the *physical storage structures*



# Relational Databases - Declarative Language

```
Database Consoles > postgres@docker > postgres_demo [postgres@docker] >
console_1 [materialize@cloud] x postgres_demo [postgres@docker] x

Database Explorer

1 CREATE TABLE IF NOT EXISTS people (
2     "id" char(21) PRIMARY KEY default nanoid(),
3     "first" char(256) NOT NULL,
4     "last" char(256) NOT NULL,
5     "yob" int8 NOT NULL
6 );
7
8 CREATE TABLE IF NOT EXISTS positions (
9     "id" char(21) PRIMARY KEY default nanoid(),
10    "name" char(256) NOT NULL,
11    "start_date" date NOT NULL,
12    "end_date" date,
13    "employee_id" char(21),
14    FOREIGN KEY (employee_id) REFERENCES people (id)
15 );
```

```
INSERT INTO people(first, last, yob)
VALUES ( first: 'amit', last: 'mor', yob: 1977),
       ( first: 'nir', last: 'pinchas', yob: 1981);

INSERT INTO positions(name, start_date, end_date, employee_id)
VALUES ( name: 'Payoneer', start_date: date('2019/3/1'), end_date: date('2020/6/1'), employee_id: '-TbRq_Mgf5Io0MDPmpX7I'),
       ( No documentation found. , end_date: null, employee_id: '-TbRq_Mgf5Io0MDPmpX7I');
```

Positions

id (pk)	name	start_date	end_date	employee_id(fk)
a2isa2n	Payoneer	1/2/2019	1/6/2020	t2a9bc
t3at78	Payoneer	1/4/2019	1/10/2020	dlli89
k8aa6d	At-Bay	7/6/2020	null	t2a9bc

People

id (pk)	first	last	yob
t2a9bc	amit	mor	1977
dlli89	nir	pinchas	1981

Relation/

# Relational Databases - JOINS

```
SELECT positions.id AS id,  
       positions.name AS company_name,  
       people.first AS employee_name,  
       people.last AS employeeed_last_name  
FROM positions  
LEFT OUTER JOIN people 1..n<->0..1: ON  
       positions.employee_id = people.id;
```

	id	company_name	employee_name	employeeed_last_name
1	KbIW0lEk0cIlcz0q8KVxH	Payoneer	amit	mor
2	4yj8Wu2v90m00Qd3Rc0Xt	At-Bay	amit	mor

Positions

id (pk)	name	start_date	end_date	employee_id (fk)
a2isa2n	Payoneer	1/2/2019	1/6/2020	t2a9bc
t3at78	Payoneer	1/4/2019	1/10/2020	dlli89
k8aa6d	At-Bay	7/6/2020	null	t2a9bc

People

id (pk)	first	last	yob
t2a9bc	amit	mor	1977
dlli89	nir	pinchas	1981

Relation/

# Relational Databases - Transactions & A.C.I.D

ACID = Atomicity, Consistency, Isolation, Durability

**Atomicity** - a guarantee that either all of the transaction succeeds or none of it does. "all or nothing".

**Consistency** - All data will be valid according to all defined rules, including any constraints, cascades, and triggers that have been applied on the database.

**Isolation** - a transaction cannot read data from any other transaction that has not yet completed.

**Durability** - once a transaction is committed, it will remain in the system – even if there's a system crash immediately following the transaction

# Relational Databases - In Practice

- You'd rarely find a company using "raw" SQL queries as it is considered error prone
- Most likely a company would use an ORM framework of some sorts
- ORM - Object (to) Relational Mapping - Libraries that abstract away the SQL expressions and allow to use a more programmatic access to data
  - Python - Django, SQLAlchemy
  - Java - Hibernate
- Python Django (similar to the left join demonstrated above):

```
Positions.objects.filter(people__id='KbIW01..')
```

# When the Relational Model Breaks

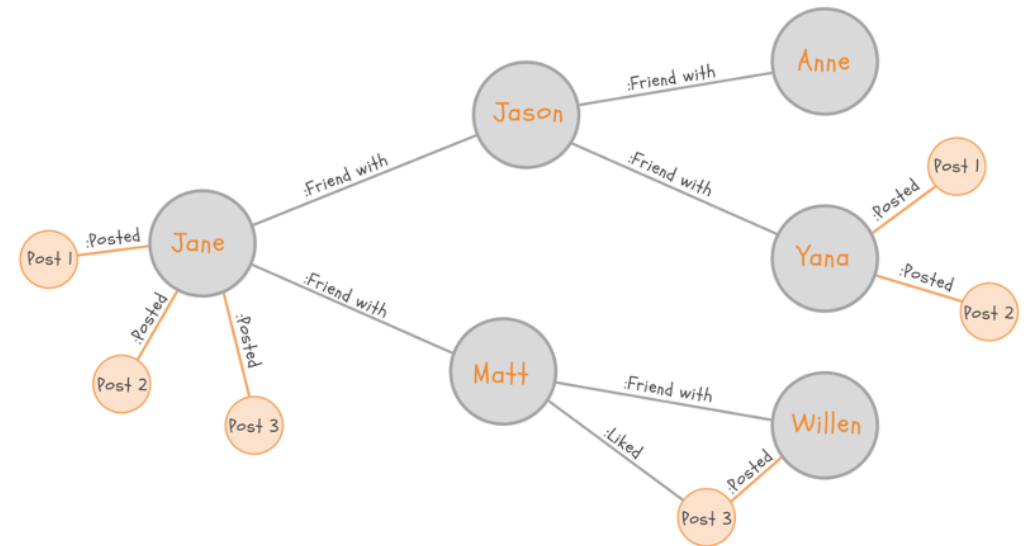
## Questions

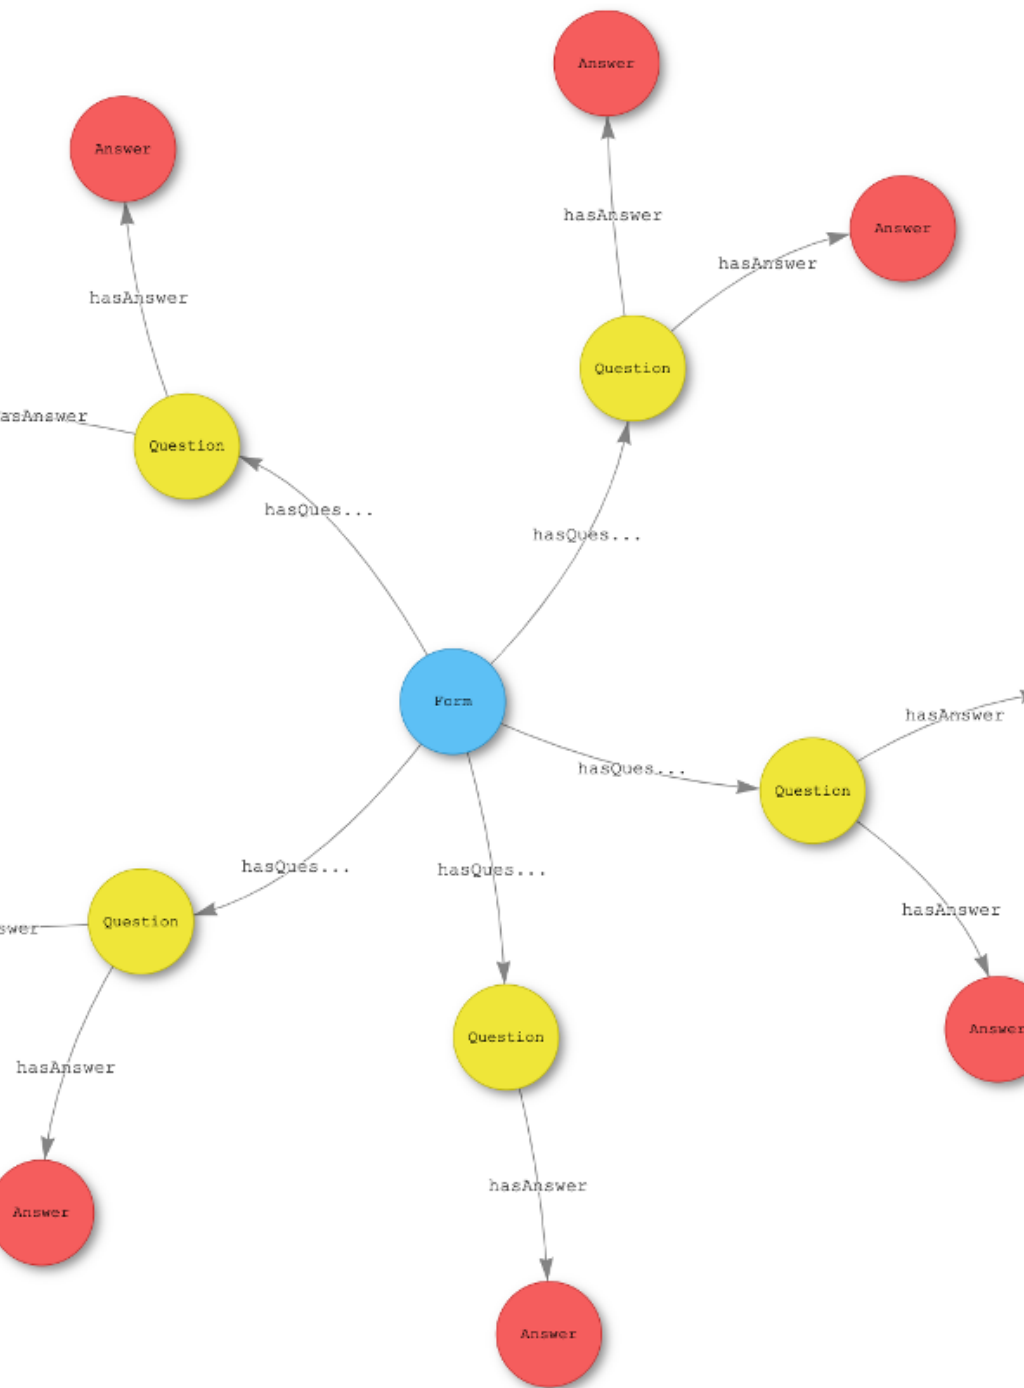
id (pk)	question	previous	next
a21sazn	what is your name?	null	b3at78
b3at78	what is your age?	a21sazn	k8aa6d
k8aa6d	have you ever committed a crime?	b3at78	null

Linked list in a relational database

# Graph Databases

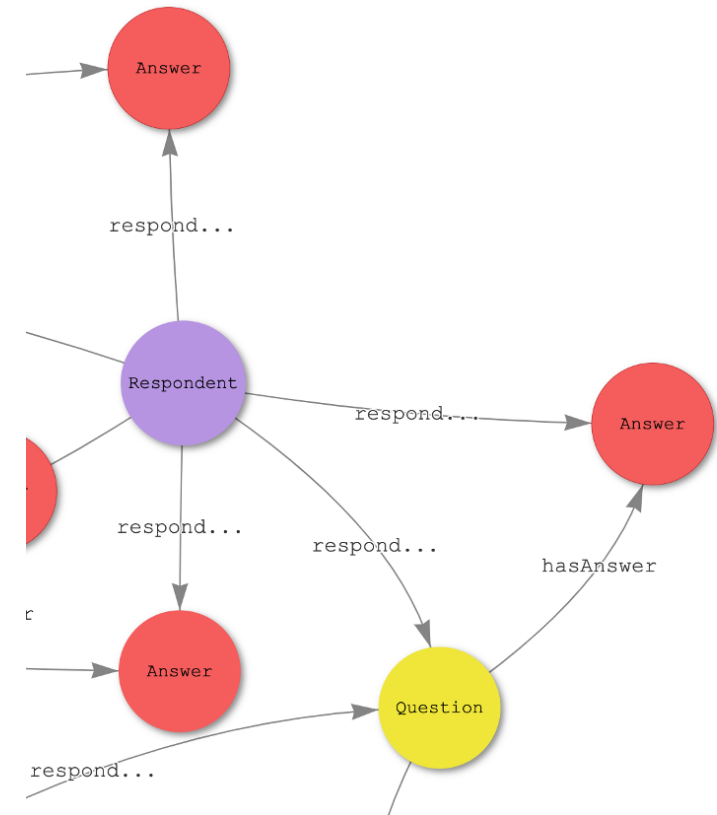
- Databases that uses **graph structures** to store data. The graph is build by nodes and edges (relationship).
- Common query languages are: Cypher, Gremlin and SparQL.
- Offers ACID transaction guarantees
- Schema-less
- Excel at recursive, graph data structures






## Modeling Questions on a Graph Database

# Modeling a Respondent





**That's a  
Questionnaire**



# At-Bay

- An insure-tech startup (5 yo)
- Raised > 200M\$; Evaluated > 1.3B\$
- x3 people in 2 years, 76 people in tlv
- working hybrid, people from north, south and center
- very diverse and heterogeneous skills are around
- best phase to do interesting things (got the money, customers and scaling the tech)
- trying to do good
- looking for back/front/full-end devs, cyber researchers, data/ml engineers

# Literature

- wikipedia
- <http://users.ece.utexas.edu/~perry/work/papers/swa-sen.pdf>
- [https://www.goodreads.com/book/show/296981.Object\\_Oriented\\_Software\\_Engineering](https://www.goodreads.com/book/show/296981.Object_Oriented_Software_Engineering)
- <https://thevaluable.dev/fighting-software-entropy/>
- E.F. Codd paper